

Linux Capabilities

Network Namespaces Commands:

- Creating and Listing Network Namespaces: `ip netns add <new namespace name>`
- Network namespace verification: `ip netns list`
- Assign Interfaces to Network Namespaces: `ip link add <veth name> type veth peer name <other veth name>`
- Veth verification: `ip link list`
- Connect interface to namespace: `ip link set <veth name> netns <namespace name>`
- Configure interface in network namespaces: `ip netns exec <network namespace> <command to run against that namespace>`

Example: `ip netns exec blue ip addr add 10.1.1.1/24 dev veth1`

More information on these commands here:

<https://blog.scottlowe.org/2013/09/04/introducing-linux-network-namespaces/>

Privileged Containers Commands in Docker:

- Check if container is in privileged mode: `docker inspect --format='{{.HostConfig.Privileged}}' [container_id]`
- Run Docker in privileged mode: `sudo docker run --privileged [image_name]`
- Run an ubuntu Container: `sudo docker run -it --privileged ubuntu`
- Test whether container has access to host: `mount -t tmpfs none /mnt`
- List disk space: `df -h`

Containers:

To add or remove Linux capabilities for a Container, include the `capabilities` field in the `securityContext` section of the Container manifest.

Yaml file example of adding capabilities:

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo-4
spec:
  containers:
    - name: sec-ctx-4
      image: gcr.io/google-samples/node-hello:1.0
      securityContext:
        capabilities:
          add: ["NET_ADMIN", "SYS_TIME"]
```

Yaml File Example of dropping capabilities:

spec:

containers:

- image: mateobur/flask

name: flask-cap

securityContext:

capabilities:

drop:

- NET_RAW

- CHOWN

Default Containers Capabilities:

- **CHOWN** - Make arbitrary changes to file UIDs and GIDs
- **DAC_OVERRIDE** - Discretionary access control (DAC) - Bypass file read, write, and execute permission checks.
- **FSETID** - Don't clear set-user-ID and set-group-ID mode bits when a file is modified; set the set-group-ID bit for a file whose GID does not match the file system or any of the supplementary GIDs of the calling process.
- **FOWNER** - Bypass permission checks on operations that normally require the file system UID of the process to match the UID of the file, excluding those operations covered by **CAP_DAC_OVERRIDE** and **CAP_DAC_READ_SEARCH**.
- **MKNOD** - Create special files using `mknod(2)`.
- **NET_RAW** - Use RAW and PACKET sockets; bind to any address for transparent proxying.
- **SETGID** - Make arbitrary manipulations of process GIDs and supplementary GID list; forge GID when passing socket credentials via UNIX domain sockets; write a group ID mapping in a user namespace.
- **SETUID** - Make arbitrary manipulations of process UIDs; forge UID when passing socket credentials via UNIX domain sockets; write a user ID mapping in a user namespace.
- **SETFCAP** - Set file capabilities.
- **SETPCAP** - If file capabilities are not supported: grant or remove any capability in the caller's permitted capability set to or from any other process.
- **NET_BIND_SERVICE** - Bind a socket to Internet domain privileged ports (port numbers less than 1024).
- **SYS_CHROOT** - Use `chroot(2)` to change to a different root directory.
- **KILL** - Bypass permission checks for sending signals. This includes use of the `ioctl(2)` `KDSIGACCEPT` operation.
- **AUDIT_WRITE** - Write records to kernel auditing log.

Proc Mount Types:

- The `/proc` directory contains virtual files that are windows into the current state of the running Linux kernel.

SELinux:

To assign SELinux labels to a Container, include the `seLinuxOptions` field in the `securityContext` section of your Pod or Container manifest. The `seLinuxOptions` field is an `SELinuxOptions` object. Here's an example that applies an SELinux level:

Yaml File Example:

```
...
securityContext:
  seLinuxOptions:
    level: "s0:c123,c456"
```